



# Dueling DQN


Paper Link: [Dueling Network Architectures for Deep Reinforcement Learning](#)

## Key Features

- The Dueling Network Architecture proposes a network architecture better suited for model-free RL than the existing standard neural network.
- The Dueling network architecture estimates the state-action value function  $Q$  using two streams consisting the state value and action advantage functions while having one common CNN network.
- The Dueling architecture can be combined with various existing and future model free RL algorithms.

## Background

This paper proposes a novel neural network architecture suitable for reinforcement learning instead of the standard neural network architecture that uses a single stream of neural networks. The Dueling network architecture as shown in Figure1 consists of two streams that represent the state value and action advantage functions, while sharing a common convolutional feature learning module.



*Figure 1. A popular single stream  $Q$ -network (**top**) and the dueling  $Q$ -network (**bottom**). The dueling network has two streams to separately estimate (scalar) state-value and the advantages for each action; the green output module implements equation (9) to combine them. Both networks output  $Q$ -values for each action.*

For an agent behaving according to a stochastic policy  $\pi$ , the values of the state-action pair  $(s, a)$  and the state  $s$  are defined as follows:

$$Q^\pi(s, a) = \mathbb{E}[R_t | s_t = s, a_t = a, \pi],$$

$$V^\pi(s) = \mathbb{E}_{a \sim \pi(s)}[Q^\pi(s, a)]$$

The paper defines another important quantity, the advantage function, relating the value and  $Q$  functions:  $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$ .

Note that  $\mathbb{E}_{a \sim \pi(s)}[A^\pi(s, a)] = 0$ . Intuitively, the state value function  $V$  measures the value of a particular state  $s$ . The  $Q$  function, however, measures the value of the particular state and action. The advantage function subtracts the value of the state from the  $Q$  function to obtain a relative measure of the importance of each action.

# Method

## The Dueling Network Architecture

From the expressions for advantage  $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$  and state-value  $V^\pi(s) = \mathbb{E}_{a \sim \pi(s)}[Q^\pi(s, a)]$ , it follows that  $\mathbb{E}_{a \sim \pi(s)}[A^\pi(s, a)] = 0$ .

Moreover, for a deterministic policy,  $a^* = \arg \max_{a' \in A} Q(s, a')$ , it follows that  $Q(s, a^*) = V(s)$  and hence  $A(s, a^*) = 0$ .

According to the above, the network in Figure1 can be expressed as an equation as follows.

$$Q(s, a; \theta, \alpha, \beta) = V(s, \theta, \beta) + A(s, a; \theta, \alpha)$$

However, the above equation has two problems. First, the  $Q(s, a; \theta, \alpha, \beta)$  is only a parameterized estimate of the true  $Q$ -function. Second, given  $Q$  is unidentifiable in that it cannot uniquely recover  $V$  and  $A$ .

To solve the above problem, when the optimal action  $a^* = \arg \max_{a' \in A} Q(s, a'; \theta, \alpha, \beta)$  is chosen,  $Q(s, a^*; \theta, \alpha, \beta)$  is made equal to  $V(s; \theta, \beta)$ . This allows the advantage function estimator can force the advantage to have zero advantage at the chosen action. This method can be Equation as

$$Q(s, a; \theta, \alpha, \beta) = V(s; \alpha, \beta) + \left( A(s, a; \theta, \alpha) - \max_{a' \in |A|} A(s, a'; \theta, \alpha) \right)$$

With this, the stream  $V(s; \theta, \beta)$  provides an estimate of the value function, while the other stream produces an estimate of the advantage function.

An alternative module replaces the max operator instead of an average:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \alpha, \beta) + \left( A(s, a; \theta, \alpha) - \frac{1}{|A|} \sum_{a'} A(s, a'; \theta, \alpha) \right)$$

The above equation loses its original meaning because  $V$  and  $A$  are now off-target by constants, but on the other hand it increases the stability of the optimization.

## Algorithm

Dueling DQN algorithm uses the algorithm of [Double DQN](#) as it is, and only the network is replaced by the Dueling network architecture.

---

### Algorithm 1: Double DQN Algorithm.

---

**input** :  $\mathcal{D}$  – empty replay buffer;  $\theta$  – initial network parameters,  $\theta^-$  – copy of  $\theta$   
**input** :  $N_r$  – replay buffer maximum size;  $N_b$  – training batch size;  $N^-$  – target network replacement freq.  
**for** episode  $e \in \{1, 2, \dots, M\}$  **do**  
    Initialize frame sequence  $\mathbf{x} \leftarrow ()$   
    **for**  $t \in \{0, 1, \dots\}$  **do**  
        Set state  $s \leftarrow \mathbf{x}$ , sample action  $a \sim \pi_{\mathcal{B}}$   
        Sample next frame  $x^t$  from environment  $\mathcal{E}$  given  $(s, a)$  and receive reward  $r$ , and append  $x^t$  to  $\mathbf{x}$   
        **if**  $|\mathbf{x}| > N_r$  **then** delete oldest frame  $x_{t_{min}}$  from  $\mathbf{x}$  **end**  
        Set  $s' \leftarrow \mathbf{x}$ , and add transition tuple  $(s, a, r, s')$  to  $\mathcal{D}$ ,  
        replacing the oldest tuple if  $|\mathcal{D}| \geq N_r$   
        Sample a minibatch of  $N_b$  tuples  $(s, a, r, s') \sim \text{Unif}(\mathcal{D})$   
        Construct target values, one for each of the  $N_b$  tuples:  
        Define  $a^{\max}(s'; \theta) = \arg \max_{a'} Q(s', a'; \theta)$   
         $y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma Q(s', a^{\max}(s'; \theta); \theta^-) & \text{otherwise.} \end{cases}$   
        Do a gradient descent step with loss  $\|y_j - Q(s, a; \theta)\|^2$   
        Replace target parameters  $\theta^- \leftarrow \theta$  every  $N^-$  steps  
    **end**  
**end**

---

## Implementation on JORLDY

- In JORLDY, only the network is replaced with the DQN algorithm.
- [Dueling DQN Agent Implementation](#)

```
from .dqn import DQN

class Dueling(DQN):
    def __init__(self, *args, **kwargs):
        if "network" in kwargs.keys():
            kwargs["network"] = "dueling"
        assert kwargs["network"] == "dueling"
        super(Dueling, self).__init__(*args, **kwargs)
```

- Dueling DQN Network Implementation

```
class Dueling(BaseNetwork):
    def __init__(self, D_in, D_out, D_hidden=512, head="mlp"):
        D_head_out = super(Dueling, self).__init__(D_in, D_hidden, head)

        self.l1_a = torch.nn.Linear(D_head_out, D_hidden)
        self.l1_v = torch.nn.Linear(D_head_out, D_hidden)

        self.l2_a = torch.nn.Linear(D_hidden, D_out)
        self.l2_v = torch.nn.Linear(D_hidden, 1)

        orthogonal_init([self.l1_a, self.l1_v])
        orthogonal_init([self.l2_a, self.l2_v], "linear")

    def forward(self, x):
        x = super(Dueling, self).forward(x)

        x_a = F.relu(self.l1_a(x))
        x_v = F.relu(self.l1_v(x))

        # A stream : action advantage
        x_a = self.l2_a(x_a) # [bs, num_action]
        x_a -= x_a.mean(dim=1, keepdim=True) # [bs, num_action]

        # V stream : state value
        x_v = self.l2_v(x_v) # [bs, 1]

        out = x_a + x_v # [bs, num_action]
        return out
```

## References

### Relevant papers

- Multi-Player Residual Advantage Learning With General Function Approximation
- Deep Reinforcement Learning with Double Q-learning